

PROBLEM STATEMENT:HOW BESTFIT THE GIVEN DATASET IS

1.DATA COLLECTION

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df=pd.read_csv(r"C:\Users\pavan\Downloads\insurance.csv")
df
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

2.DATA CLEANING & PREPROCESSING

```
In [3]: df.head()
```

```
Out[3]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [4]: df.tail()
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
In [5]: df.shape
```

```
Out[5]: (1338, 7)
```

```
In [6]: df.describe()
```

```
Out[6]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [8]: df.isna().sum()
```

```
Out[8]: age          0  
sex          0  
bmi          0  
children     0  
smoker       0  
region       0  
charges      0  
dtype: int64
```

```
In [9]: df['region'].value_counts()
```

```
Out[9]: region  
southeast    364  
southwest    325  
northwest    325  
northeast    324  
Name: count, dtype: int64
```

```
In [10]: convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```


Out[10]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	0	33.770	1	no	southeast	1725.55230
2	28	0	33.000	3	no	southeast	4449.46200
3	33	0	22.705	0	no	northwest	21984.47061
4	32	0	28.880	0	no	northwest	3866.85520
...
1333	50	0	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

3.DATA VISUALIZATION

```
In [11]: sns.lmplot(x='bmi',y='charges',order=2,data=df,ci=None)  
plt.show()
```




```
In [12]: x=np.array(df['bmi']).reshape(-1,1)  
y=x=np.array(df['charges']).reshape(-1,1)
```


```
In [13]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=0)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

1.0

```
In [14]: y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```




```
In [15]: df700=df[:][:700]  
sns.lmplot(x='bmi',y='charges',order=2,ci=None,data=df700)  
plt.show()
```



```
In [16]: df700.fillna(method='ffill',inplace=True)
```

```
In [17]: x=np.array(df700["bmi"]).reshape(-1,1)
y=np.array(df700['charges']).reshape(-1,1)
```


```
In [18]: df700.dropna(inplace=True)
```

```
In [19]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

```
0.020508709612662157
```

In [20]:

```
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



EVALUATION

```
In [21]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [22]: lr=LinearRegression()
lr.fit(x_train,y_train)
y_pred=lr.predict(x_test)
r2=r2_score(y_test,y_pred)
print(r2)
```

0.020508709612662157

accuracy is 0.03442

RIDGE REGRESSION

Type *Markdown* and LaTeX: α^2

```
In [23]: convert={'smoker':{'yes':1,"no":0}}
df=df.replace(convert)
df
```

Out[23]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	0	33.770	1	0	southeast	1725.55230
2	28	0	33.000	3	0	southeast	4449.46200
3	33	0	22.705	0	0	northwest	21984.47061
4	32	0	28.880	0	0	northwest	3866.85520
...
1333	50	0	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [24]: convert={'region':{'southeast':1,"southwest":2,"northeast":3,"northwest":4}}
df=df.replace(convert)
df
```

Out[24]:

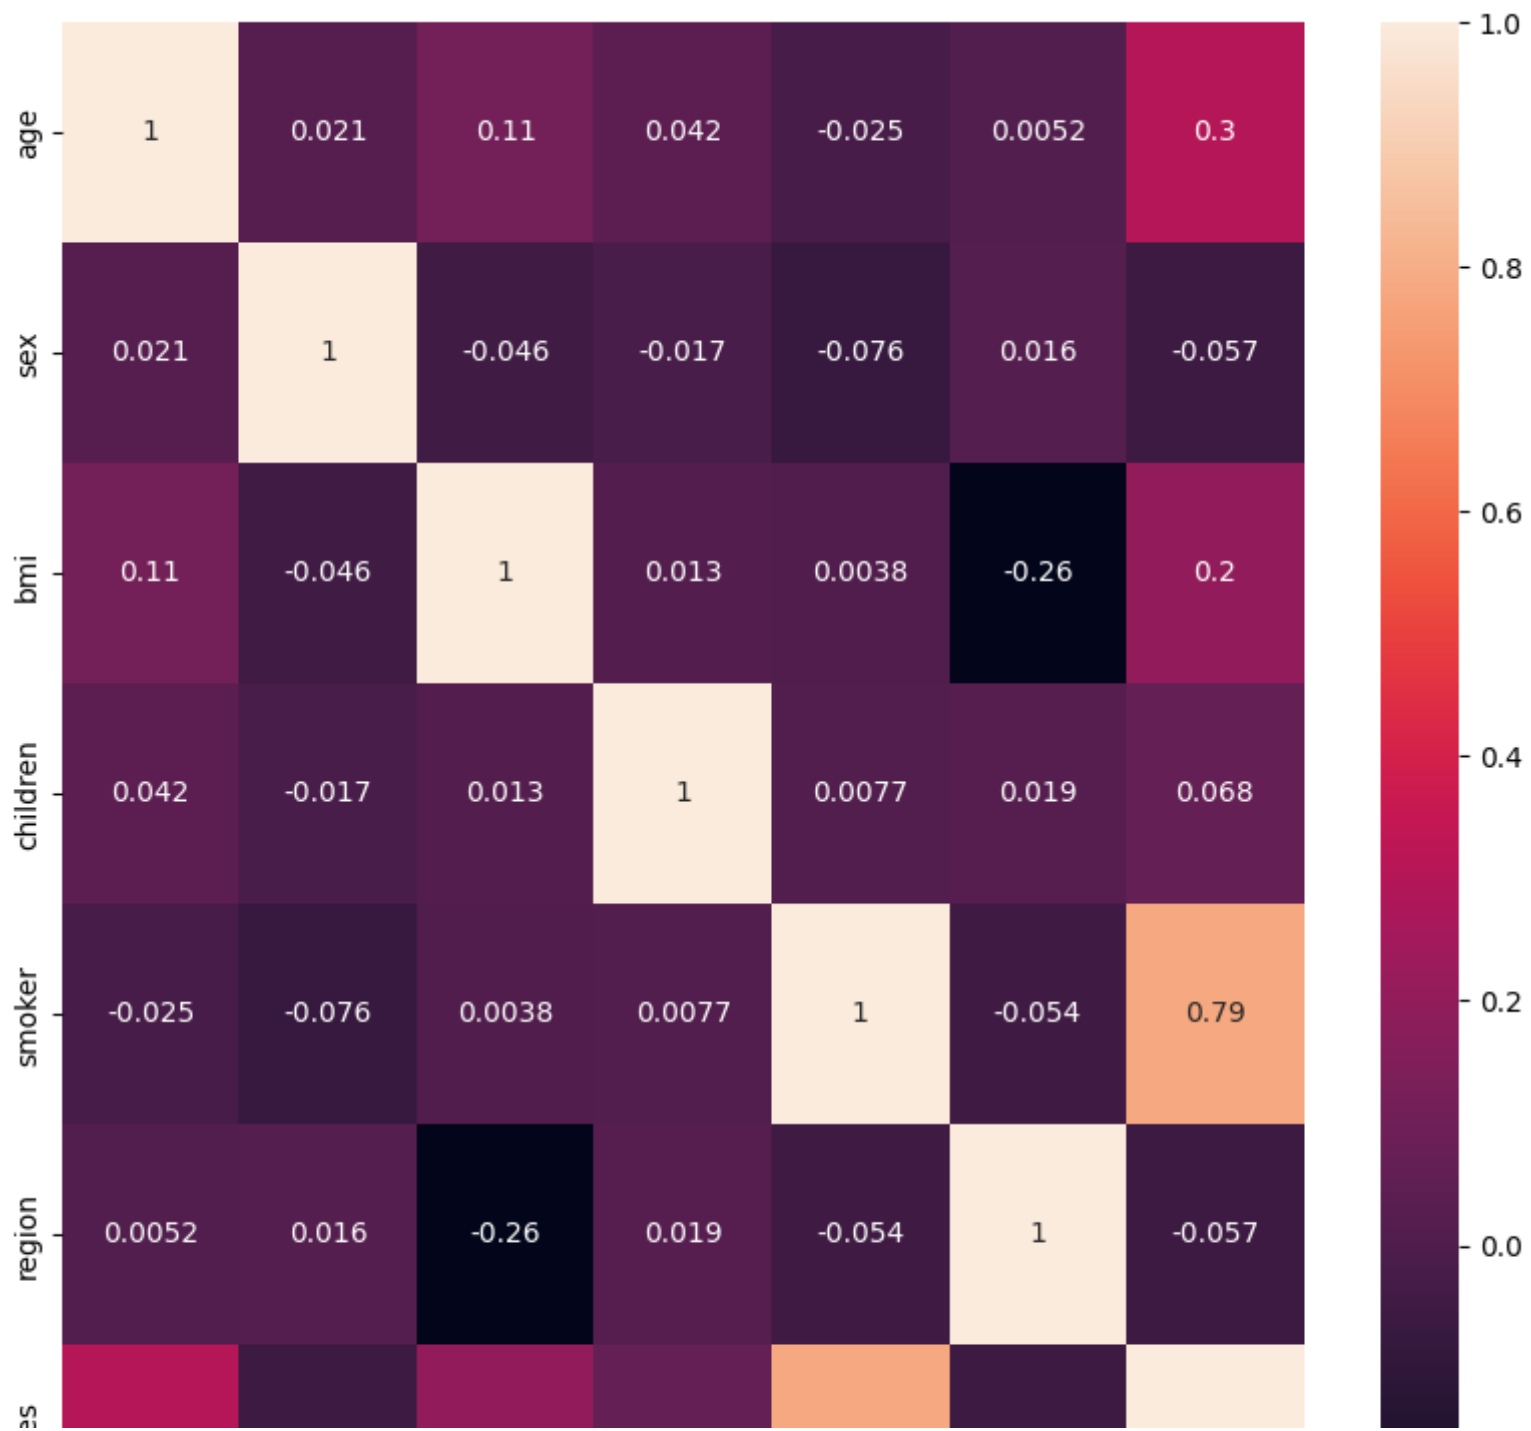
	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	2	16884.92400
1	18	0	33.770	1	0	1	1725.55230
2	28	0	33.000	3	0	1	4449.46200
3	33	0	22.705	0	0	4	21984.47061
4	32	0	28.880	0	0	4	3866.85520
...
1333	50	0	30.970	3	0	4	10600.54830
1334	18	1	31.920	0	0	3	2205.98080
1335	18	1	36.850	0	0	1	1629.83350
1336	21	1	25.800	0	0	2	2007.94500
1337	61	1	29.070	0	1	4	29141.36030


1338 rows × 7 columns

```
In [25]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [26]: plt.figure(figsize = (10, 10))  
sns.heatmap(df.corr(), annot = True)
```

```
Out[26]: <Axes: >
```



```
In [27]: features=df.columns[0:1]
target=df.columns[-1]
```

```
In [28]: x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=1)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
```

The dimension of X_train is (936, 1)
The dimension of X_test is (402, 1)

```
In [29]: lr = LinearRegression()
#Fit model
lr.fit(x_train, y_train)
#predict
actual = y_test
train_score_lr = lr.score(x_train, y_train)
test_score_lr = lr.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.0910963973805714
The test score for lr model is 0.08490473916580776

```
In [30]: ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.09109639711159634


The test score for ridge model is 0.08490538609860176

```
In [31]: plt.figure(figsize=(10,10))
```

Out[31]: <Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

```
In [32]: plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge;\alpha=0.7')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,color='green',label='LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



LASSO REGRESSION

```
In [33]: lasso= Lasso(alpha=10)
lasso.fit(x_train,y_train)
#train and test scorefor ridge regression
train_score_ls = lasso.score(x_train, y_train)
test_score_ls= lasso.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for lasso model is {}".format(train_score_ls))
print("The test score for lasso model is {}".format(test_score_ls))
```

Ridge Model:


The train score for lasso model is 0.09109639395809055
The test score for lasso model is 0.08490704421828055

```
In [34]: plt.figure(figsize=(10,10))
```

Out[34]: <Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

```
In [35]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")  
plt.show()
```



```
In [36]: from sklearn.linear_model import LassoCV
```

```
In [37]: #using the linear cv model
from sklearn.linear_model import RidgeCV
#cross validation
ridge_cv=RidgeCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print(ridge_cv.score(x_train,y_train))
print(ridge_cv.score(x_test,y_test))

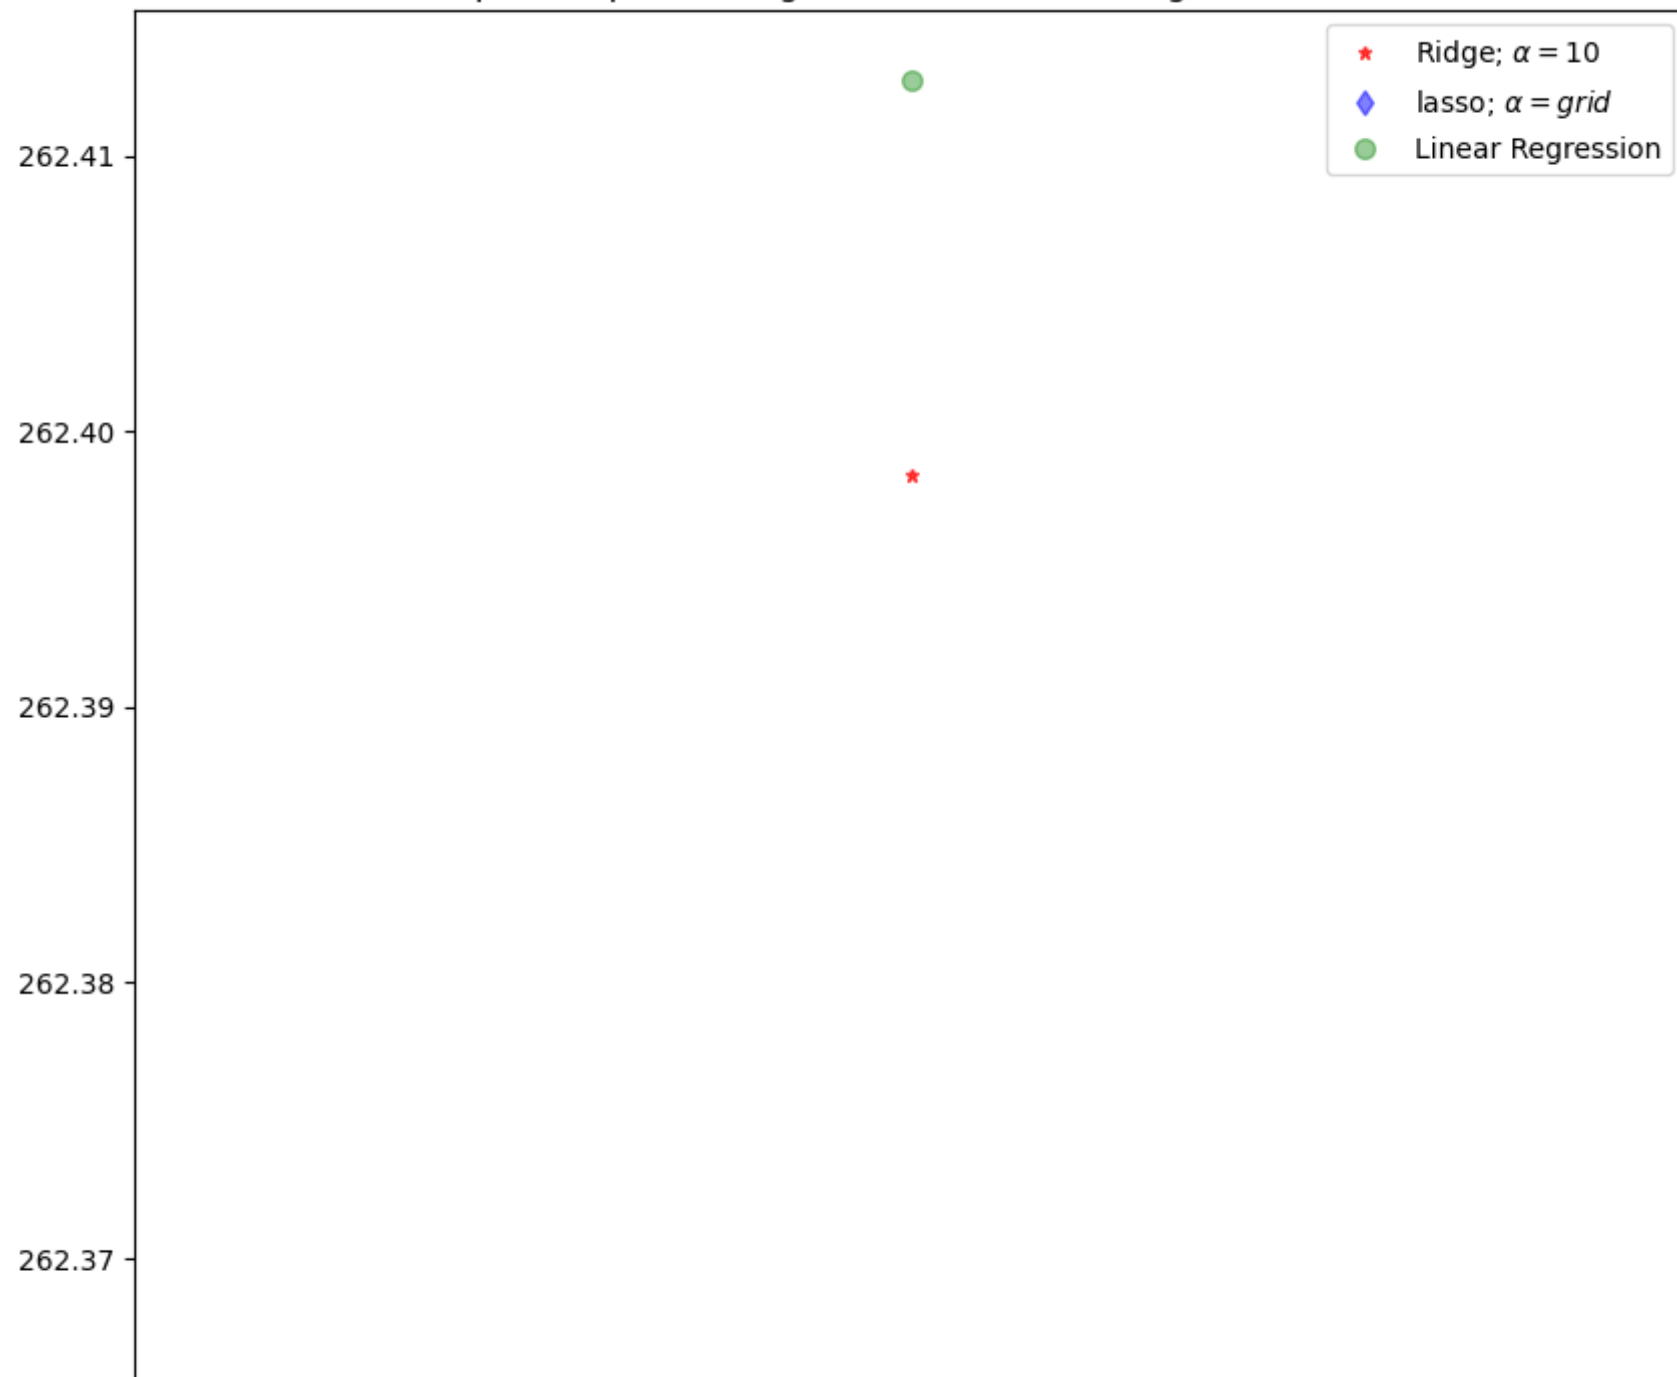
0.09109639711159612
0.08490538609884779
```


```
In [38]: #using the linear cv model
from sklearn.linear_model import LassoCV
#cross validation
lasso_cv=LassoCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))

0.09109639395809055
0.08490704421828055
```

```
In [39]: ot size
figure(figsize = (10, 10))
plot for ridge regression
plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge;  $\alpha=10$ ')
plot for lasso regression
plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\alpha = \text{grid}$ ')
plot for linear model
plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
ate axis
xticks(rotation = 90)
legend()
title("Comparison plot of Ridge, Lasso and Linear regression model")
show()
```


Comparison plot of Ridge, Lasso and Linear regression model





ElasticNet

```
In [40]: from sklearn.linear_model import ElasticNet
```

```
In [41]: el=ElasticNet()  
el.fit(x_train,y_train)  
print(el.coef_)  
print(el.intercept_)  
el.score(x,y)
```

```
[261.74450967]  
3115.083177426244
```

```
Out[41]: 0.08930616764094623
```

```
In [42]: y_pred_elastic=el.predict(x_train)
```

```
In [43]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)  
print(mean_squared_error)
```

```
135077142.70714515
```

accuracy was 0.08930

LOGISTIC REGRESSION

```
In [44]: import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [45]: df=pd.read_csv(r"C:\Users\pavan\Downloads\insurance.csv")
df
```

Out[45]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [46]: df.shape
```

```
Out[46]: (1338, 7)
```

```
In [47]: pd.set_option('display.max_rows',10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
```

```
In [48]: print('This Dataset has %d rows and %d columns'%(df.shape))
```

This Dataset has 1338 rows and 7 columns

```
In [49]: df.head()
```

```
Out[49]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [50]: df.tail()
```

```
Out[50]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
In [51]: df.describe()
```

```
Out[51]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [52]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [53]: df.isna().sum()
```

```
Out[53]: age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```

```
In [54]: convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

184	44	male	30.690	2	0	southeast	7731.427100
185	36	male	41.895	3	1	northeast	43753.337050
186	26	female	29.920	2	0	southeast	3981.976800
187	30	female	30.900	3	0	southwest	5325.651000
188	41	female	32.200	1	0	southwest	6775.961000
189	29	female	32.110	2	0	northwest	4922.915900
190	61	male	31.570	0	0	southeast	12557.605300
191	36	female	26.200	0	0	southwest	4883.866000
192	25	male	25.740	0	0	southeast	2137.653600
193	56	female	26.600	1	0	northwest	12044.342000
194	18	male	34.430	0	0	southeast	1137.469700
195	19	male	30.590	0	0	northwest	1639.563100
196	39	female	32.800	0	0	southwest	5649.715000

```
In [55]: convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

45	55	0	37.300	0	0	southwest	20630.283510
46	18	1	38.665	2	0	northeast	3393.356350
47	28	1	34.770	0	0	northwest	3556.922300
48	60	1	24.530	0	0	southeast	12629.896700
49	36	0	35.200	1	1	southeast	38709.176000
50	18	1	35.625	0	0	northeast	2211.130750
51	21	1	33.630	2	0	northwest	3579.828700
52	48	0	28.000	1	1	southwest	23568.272000
53	36	0	34.430	0	1	southeast	37742.575700
54	40	1	28.690	3	0	northwest	8059.679100
55	58	0	36.955	2	1	northwest	47496.494450
56	58	1	31.825	2	0	northeast	13607.368750
57	18	0	31.680	2	1	southeast	34303.167200


```
In [56]: convert={"region":{"southeast":1,"southwest":2,"northeast":3,"northwest":4}}
df=df.replace(convert)
df
```

64	20	1	22.420	0	1	4	14711.743600
65	19	1	28.900	0	0	2	1743.214000
66	61	1	39.100	2	0	2	14235.072000
67	40	0	26.315	1	0	4	6389.377850
68	40	1	36.190	0	0	1	5920.104100
69	28	0	23.980	3	1	1	17663.144200
70	27	1	24.750	0	1	1	16577.779500
71	31	0	28.500	5	0	3	6799.458000
72	53	1	28.100	3	0	2	11741.726000
73	58	0	32.010	1	0	1	11946.625900
74	44	0	27.400	2	0	2	7726.854000
75	57	0	34.010	0	0	4	11356.660900
76	29	1	29.590	1	0	1	3947.413100
77	21	0	25.500	0	0	1	1500.100700

```
In [57]: features_matrix=df.iloc[:,0:4]
```


```
In [58]: target_vector=df.iloc[:, -3]
```

```
In [59]: print('The Feature Matrix has %d Rows and %d columns(s)'%(features_matrix.shape))
print('The Target Matrix has %d Rows and %d columns(s)%(np.array(target_vector).reshape(-1,1).shape))
```

```
The Feature Matrix has 1338 Rows and 4 columns(s)
The Target Matrix has 1338 Rows and 1 columns(s)
```

```
In [60]: import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [61]: sns.barplot(x='smoker', y='age', data=df, color="mediumturquoise")  
plt.show()
```



```
In [62]: features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```



```
In [71]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.05)
         lo=LogisticRegression()
         lo.fit(x_train,y_train)
         print(lo.score(x_test,y_test))
```

0.835820895522388

C:\Users\pavan\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

for logistic regression the accuracy was 0.8059

DECISIONTREE

```
In [72]: import numpy as np
         import pandas as pd
         import seaborn as sns
         from sklearn.model_selection import train_test_split
         from sklearn.tree import DecisionTreeClassifier
```

```
In [73]: df=pd.read_csv(r"C:\Users\pavan\Downloads\insurance.csv")
df
```

8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800
11	62	female	26.290	0	yes	southeast	27808.725100
12	23	male	34.400	0	no	southwest	1826.843000
13	56	female	39.820	0	no	southeast	11090.717800
14	27	male	42.130	0	yes	southeast	39611.757700
15	19	male	24.600	1	no	southwest	1837.237000
16	52	female	30.780	1	no	northeast	10797.336200
17	23	male	23.845	0	no	northeast	2395.171550
18	56	male	40.300	0	no	southwest	10602.385000
19	30	male	35.300	0	yes	southwest	36837.467000
20	60	female	36.005	0	no	northeast	13228.846950

```
In [74]: df.shape
```

```
Out[74]: (1338, 7)
```

```
In [75]: df.isnull().sum()
```

```
Out[75]: age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

```
In [76]: df.isna().any()
```

```
Out[76]: age          False  
sex           False  
bmi           False  
children      False  
smoker        False  
region        False  
charges       False  
dtype: bool
```

```
In [77]: df['region'].value_counts()
```

```
Out[77]: region  
southeast    364  
southwest    325  
northwest    325  
northeast    324  
Name: count, dtype: int64
```

```
In [78]: convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

55	58	0	36.955	2	yes	northwest	47496.494450
56	58	1	31.825	2	no	northeast	13607.368750
57	18	0	31.680	2	yes	southeast	34303.167200
58	53	1	22.880	1	yes	southeast	23244.790200
59	34	1	37.335	2	no	northwest	5989.523650
60	43	0	27.360	3	no	northeast	8606.217400
61	25	0	33.660	4	no	southeast	4504.662400
62	64	0	24.700	1	no	northwest	30166.618170
63	28	1	25.935	1	no	northwest	4133.641650
64	20	1	22.420	0	yes	northwest	14711.743800
65	19	1	28.900	0	no	southwest	1743.214000
66	61	1	39.100	2	no	southwest	14235.072000
67	40	0	26.315	1	no	northwest	6389.377850

```
In [79]: convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

34	28	0	36.400	1	1	southwest	51194.559140
35	19	0	20.425	0	0	northwest	1625.433750
36	62	1	32.965	3	0	northwest	15612.193350
37	26	0	20.800	0	0	southwest	2302.300000
38	35	0	36.670	1	1	northeast	39774.276300
39	60	0	39.900	0	1	southwest	48173.361000
40	24	1	26.600	0	0	northeast	3046.062000
41	31	1	36.630	2	0	southeast	4949.758700
42	41	0	21.780	1	0	southeast	6272.477200
43	37	1	30.800	2	0	southeast	6313.759000
44	38	0	37.050	1	0	northeast	6079.671500
45	55	0	37.300	0	0	southwest	20630.283510
46	18	1	38.665	2	0	northeast	3393.356350

```
In [80]: x=["bmi","children"]
y=["Yes","No"]
all_inputs=df[x]
all_classes=df["sex"]
```

```
In [81]: (x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.03)
```

```
In [82]: clf=DecisionTreeClassifier(random_state=0)
```



```
In [83]: clf.fit(x_train,y_train)
```

```
Out[83]: 

▼ DecisionTreeClassifier



DecisionTreeClassifier(random_state=0)


```

```
In [84]: score=clf.score(x_test,y_test)  
print(score)
```

```
0.5853658536585366
```

the accuracy for decision tree was 0.4878

RANDOM FOREST

```
In [85]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt ,seaborn as sns
```

```
In [86]: df=pd.read_csv(r"C:\Users\pavan\Downloads\insurance.csv")
df
```

57	18	male	31.680	2	yes	southeast	34303.167200
58	53	female	22.880	1	yes	southeast	23244.790200
59	34	female	37.335	2	no	northwest	5989.523650
60	43	male	27.360	3	no	northeast	8606.217400
61	25	male	33.660	4	no	southeast	4504.662400
62	64	male	24.700	1	no	northwest	30166.618170
63	28	female	25.935	1	no	northwest	4133.641650
64	20	female	22.420	0	yes	northwest	14711.743800
65	19	female	28.900	0	no	southwest	1743.214000
66	61	female	39.100	2	no	southwest	14235.072000
67	40	male	26.315	1	no	northwest	6389.377850
68	40	female	36.190	0	no	southeast	5920.104100
69	28	male	23.980	3	yes	southeast	17663.144200

```
In [87]: df.shape
```

```
Out[87]: (1338, 7)
```

```
In [88]: df['region'].value_counts()
```

```
Out[88]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [89]: df['bmi'].value_counts()

28.595	6
37.100	6
31.730	6
33.000	6
27.740	6
29.830	6
25.175	6
26.410	6
27.835	5
33.660	5
28.900	5
25.080	5
28.500	5
26.695	5
23.210	5
32.395	5
31.825	5
33.155	5
29.640	5
36.850	5

```
In [90]: m={"sex":{"female":1,"male":0}}
df=df.replace(m)
print(df)
```

35	19	0	20.425	0	no	northwest	1625.433750
36	62	1	32.965	3	no	northwest	15612.193350
37	26	0	20.800	0	no	southwest	2302.300000
38	35	0	36.670	1	yes	northeast	39774.276300
39	60	0	39.900	0	yes	southwest	48173.361000
40	24	1	26.600	0	no	northeast	3046.062000
41	31	1	36.630	2	no	southeast	4949.758700
42	41	0	21.780	1	no	southeast	6272.477200
43	37	1	30.800	2	no	southeast	6313.759000
44	38	0	37.050	1	no	northeast	6079.671500
45	55	0	37.300	0	no	southwest	20630.283510
46	18	1	38.665	2	no	northeast	3393.356350
47	28	1	34.770	0	no	northwest	3556.922300
48	60	1	24.530	0	no	southeast	12629.896700
49	36	0	35.200	1	yes	southeast	38709.176000
50	18	1	35.625	0	no	northeast	2211.130750
51	21	1	33.630	2	no	northwest	3579.828700
52	48	0	28.000	1	yes	southwest	23568.272000
53	36	0	34.430	0	yes	southeast	37742.575700
54	40	1	28.690	3	no	northwest	8059.679100

```
In [91]: n={"smoker":{"yes":1,"no":0}}
df=df.replace(n)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920
10	25	0	26.220	0	0	northeast	2721.320800
11	62	1	26.290	0	1	southeast	27808.725100
12	23	0	34.400	0	0	southwest	1826.843000
13	56	1	39.820	0	0	southeast	11090.717800
14	27	0	42.130	0	1	southeast	39611.757700
15	19	0	24.600	1	0	southwest	1837.237000
16	52	1	30.780	1	0	northeast	10797.336200
17	23	0	23.845	0	0	northeast	2395.171550
18	55	0	29.300	0	0	southwest	16622.005000

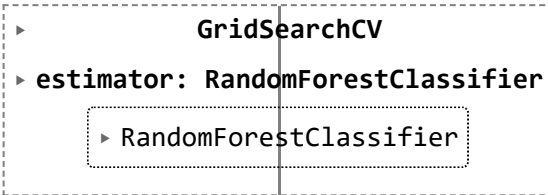
```
In [92]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[92]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [93]: rf=RandomForestClassifier()  
params={'max_depth':[2,3,5,20],  
        'min_samples_leaf':[5,10,20,50,100,200],  
        'n_estimators':[10,25,30,50,100,200]}
```

```
In [94]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

```
Out[94]:
```



```
  ▸ GridSearchCV  
  ▸ estimator: RandomForestClassifier  
    ▸ RandomForestClassifier
```

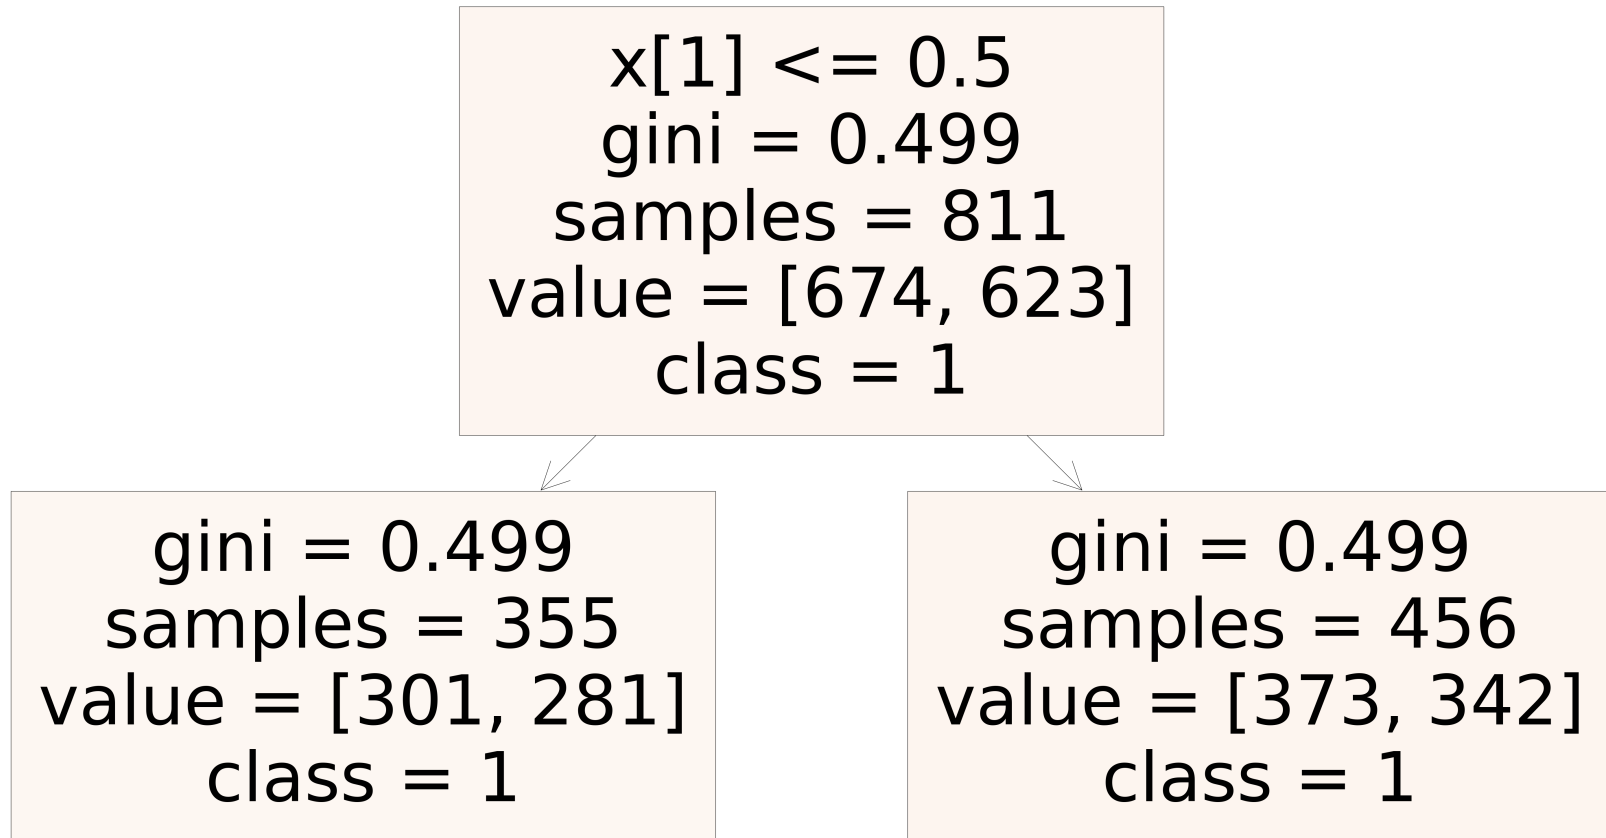
```
In [95]: grid_search.best_score_
```

```
Out[95]: 0.5227462953451654
```

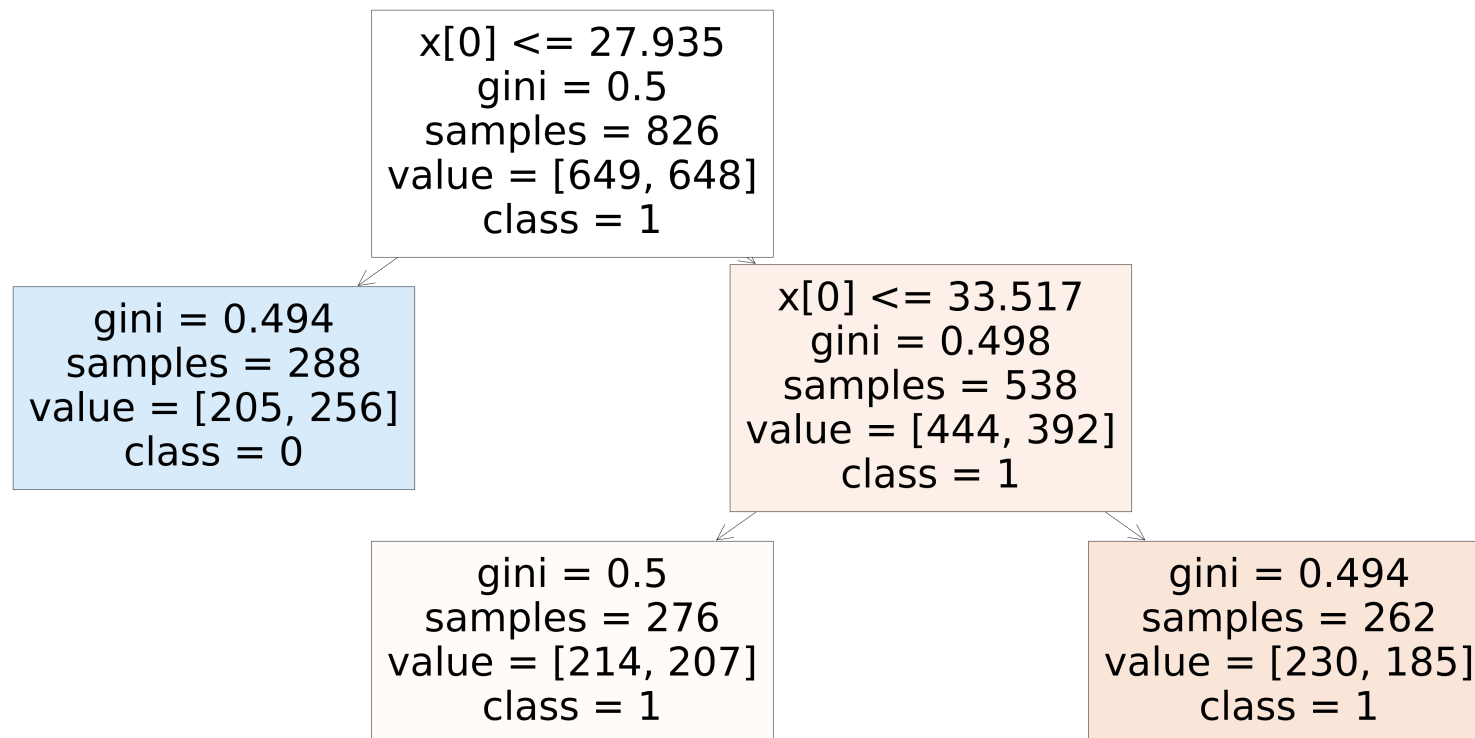
```
In [96]: rf_best=grid_search.best_estimator_  
print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=200)
```

```
In [97]: from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))  
plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```



```
In [98]: from sklearn.tree import plot_tree
plt.figure(figsize=(70,30))
plot_tree(rf_best.estimators_[6],class_names=["1","0"],filled=True);
```



```
In [99]: rf_best.feature_importances_
```

```
Out[99]: array([0.61264404, 0.38735596])
```

```
In [100]: rf=RandomForestClassifier(random_state=0)
```



```
In [101]: rf.fit(x_train,y_train)
```

```
Out[101]: RandomForestClassifier  
RandomForestClassifier(random_state=0)
```

```
In [102]: score=rf.score(x_test,y_test)  
print(score)
```

```
0.5853658536585366
```

the accuracy is 0.43902

CONCLUSION

The given dataset is "Insurance",we need to find the bestfit Model. As per the data set ,we have used different types of models,that different models got different types of accuracies.In this process Ridge and Lasso got same accuracy i.e 0.091.so,we should not consider that.For the ElasticNet model I got the accuracy of 0.089306.For the highest accuracy,I have done so many models among those ElasticNet got highest accuracy.I have done so many visuvalization graphs as per the given Features.

So,the ElasticNet Model is the bestfit for the given Dataset "Insurance".

```
In [ ]:
```